

Lab 3: Compiler Construction

Parser with Flex and Bison

Technical Report

Manmadha Rao Kuttum (2025MCS2101)
Jithendra Gannavarapu (2025MCS2743)

December 24, 2025

Abstract

This report documents the design and implementation of a compiler for a C-like imperative programming language. The project utilizes **Flex** for lexical analysis and **Bison** for parsing to generate an Abstract Syntax Tree (AST). Key features include integer arithmetic with **unary operator support**, variable declarations, nested block scopes, and structured control flow. Furthermore, the project extends beyond a basic front-end by including a **runtime interpreter** capable of executing the AST, printing output, and enforcing runtime safety checks such as **division-by-zero detection**.

1 Overview

The objective of this lab was to build a robust front-end compiler capable of verifying the syntax and semantics of a custom programming language. The compiler reads source code from standard input or files, verifies lexical and syntactic correctness, and constructs a hierarchical AST representation.

The project repository is maintained at:

<https://github.com/JithendraGannavarapu/CornerStoneProject/tree/main/lab3>

2 Description of the Designed Language

The language is designed to mimic the fundamental structure of C, focusing on integer arithmetic and structured control flow.

2.1 Key Features

- **Data Types:** Strictly Integers (literals and variables).
- **Declarations:** Explicit typing using the `var` keyword (e.g., `var x = 10;`).
- **Control Flow:**
 - `if (condition) { ... } else { ... }`
 - `while (condition) { ... }`

- **Scoping:** Nested block scopes { ... } are supported, allowing for complex, recursive structures.
- **Comments:**
 - Single-line: // (Matches until newline).
 - Multi-line: /* ... */ (Matches across line breaks using the regex pattern: "/*([*]|(n*\+[*]))*n*\+/\").

3 Internal Structure and Workflow

The compiler functions as a two-stage pipeline where the **Lexer** (Lexical Analyzer) and **Parser** (Syntax Analyzer) work in tandem to transform raw text into a structured Abstract Syntax Tree (AST). This section details the internal mechanics of this interaction.

3.1 Data Flow Architecture

The interaction between the Lexer and Parser is driven by the pull-model architecture of `yyparse()`.

1. **Controller:** The Parser (`yyparse()`) controls the flow. It requests the next token only when it needs one to complete a grammar rule.
2. **Provider:** The Lexer (`yylex()`) scans the input stream and returns a single token ID (e.g., INTEGER, VAR).
3. **Data Channel:** The actual content of the token (like the value 10 or name "x") is passed through a global union variable called `yylval`.

3.2 Internal Code Structure

3.2.1 1. The Global Union (`yylval`)

To allow the Lexer to pass different data types (int, string, or AST Node) to the Parser, we defined a C union in Bison. This is the critical "bridge" between the two components.

```

1 %union {
2     int ival;          /* For integer literals (e.g., 10) */
3     char* sval;        /* For identifiers (e.g., "count") */
4     struct ASTNode* nval; /* For AST Nodes (expressions, statements) */
5 }
```

Listing 1: The Data Bridge (Bison Definition)

3.2.2 2. The Lexer's Role (`lexer.1`)

The Lexer calculates the value and places it into the bridge before returning the token type.

```

1 [0-9]+ {
2     yylval.ival = atoi(yytext); /* 1. Convert text to integer */
3     return INTEGER;           /* 2. Notify Parser we found an INTEGER */
4 }
5
6 [a-zA-Z_][a-zA-Z0-9_]* {
7     yylval.sval = strdup(yytext); /* 1. Copy string to heap */
```

```

8     return IDENTIFIER;           /* 2. Notify Parser we found an ID */
9 }

```

Listing 2: Lexer Populating the Bridge

3.2.3 3. The Parser's Role (parser.y)

The Parser reads from the bridge using the \$ notation. \$1 refers to the first component's value, \$2 the second, and so on.

```

1 /* Rule: term -> term PLUS factor */
2 term: term PLUS factor {
3     /* $$ = Output Node
4      $1 = Left Child (term)
5      $3 = Right Child (factor) */
6     $$ = create_binop("+", $1, $3);
7 }

```

Listing 3: Parser Consuming Data

3.3 Execution Lifecycle

The internal workflow for processing a single statement like `var x = 10;` proceeds as follows:

1. **Initialization:** `main()` calls `yyparse()`.
2. **Token Request 1:** `yyparse` calls `yylex`. Lexer finds `var`, returns `VAR`.
3. **Token Request 2:** `yyparse` calls `yylex`. Lexer finds `x`.
 - Action: Sets `yyval.sval = "x"`.
 - Return: `IDENTIFIER`.
4. **Token Request 3:** `yyparse` calls `yylex`. Lexer finds `=`, returns `ASSIGN`.
5. **Token Request 4:** `yyparse` calls `yylex`. Lexer finds `10`.
 - Action: Sets `yyval.ival = 10`.
 - Return: `INTEGER`.
6. **Reduction:** The Parser sees the pattern `VAR IDENTIFIER ASSIGN INTEGER` matches a grammar rule. It executes the semantic action to create a `VAR_DECL` AST node using the data in `yyval`.

4 Delimiters and File Structure

In Flex (Lexer) and Bison (Parser) files, specific delimiters are used to switch between **C Code**, **Regular Expressions**, and **Grammar Rules**. These markers separate the distinct sections of the code.

4.1 The C Block Delimiter: %{ ... %}

This delimiter is used in the **Definitions Section** at the very top of the file.

- **Meaning:** "Everything inside here is raw C code. Copy it directly to the generated file."
- **Usage:** It is primarily used for `#include` statements and global variable declarations.

```
1 %{  
2     #include <stdio.h>  
3     #include "ast.h"    /* Include C headers here */  
4 %}
```

Listing 4: Usage of C Block Delimiter

4.2 The Section Separator: %%

This is the primary delimiter that divides the file into three distinct sections.

1. **Definitions Section:** Imports and token definitions.
 2. **Rules Section:** Regular expressions (Lexer) or Grammar rules (Parser).
 3. **User Code Section:** C functions like `main()` or helpers.
- The **first %%** marks the end of definitions and the start of rules.
 - The **second %%** marks the end of rules and the start of raw C code.

4.3 Bison Declarations

In addition to the block delimiters, specific % keywords are used for declarations in the parser:

- `%union { ... }`: Defines the semantic value types (int, string, AST node).
- `%token`: Defines terminal symbols coming from the Lexer.
- `%type`: Defines the return type of a non-terminal rule.
- `%left / %right`: Defines operator precedence and associativity.

4.4 Structure Visualization

The following snippet demonstrates how these delimiters structure a real file:

```
1 /* SECTION 1: DEFINITIONS */  
2 %{  
3     #include <stdio.h> /* C Code Block */  
4 %}  
5  
6 %token INTEGER          /* Bison Declaration */  
7  
8 %%                      /* SEPARATOR: Switch to Rules */  
9  
10 /* SECTION 2: RULES */  
11 [0-9]+ { return INTEGER; }  
12  
13 %%                      /* SEPARATOR: Switch to User Code */
```

```

14
15 /* SECTION 3: USER CODE */
16 int main() {
17     yylex();
18 }
```

Listing 5: File Structure with Delimiters

5 Lexical Analysis (The Lexer)

The Lexer (`lexer.y`) acts as the "front desk" of the compiler. It transforms the raw stream of input characters into meaningful **Tokens**.

5.1 Tokenization Strategy

We utilized **Flex** to implement the tokenizer. It uses Regular Expressions (Regex) to match patterns and return specific Token IDs to the parser.

- **Keywords:** `var, if, else, while`.
- **Identifiers:** `[a-zA-Z_][a-zA-Z0-9_]*` (Standard C naming conventions).
- **Integers:** `[0-9]+`.
- **Operators:** `+, -, *, /, ==, <=`, etc.

5.2 Design Decision: The `yyval` Union

A critical implementation detail is the data transfer between Lexer and Parser. When the Lexer identifies a value, it must pass the *content* (not just the type) to the Parser. We used the `%union` construct in Bison to define `yyval`:

- `ival (int)`: Stores the numeric value of INTEGER tokens.
- `sval (char*)`: Stores the name of IDENTIFIER tokens.
- `nval (ASTNode*)`: Used by the parser for non-terminals (e.g., `expression`).

5.3 Error Handling: Panic Mode

To ensure the compiler does not produce invalid executables from garbage input, we implemented a **Panic Mode** strategy.

- **Logic:** A catch-all rule `(.)` at the end of the lexer file detects any unknown character.
- **Action:** It prints an error to `stderr` and calls `exit(1)` immediately. This ensures the operating system receives a failure exit code.

6 Grammar Design & Parsing (The Parser)

The Parser (`parser.y`), generated by **Bison**, defines the grammatical rules and builds the AST.

6.1 Recursive Grammar & Robustness

The grammar is designed recursively to support infinite nesting.

```
statement -> block -> statement_list -> statement
```

This design allows for constructs like `while` loops inside `if` blocks inside other `while` loops. We verified this using a "Stress Test" containing 5 levels of nesting.

6.2 Operator Precedence

To respect standard mathematical order of operations (PEMDAS), precedence is encoded directly into the grammar hierarchy. We introduced a specific `unary` rule to handle negative numbers correctly:

1. `expression` (Lowest: `==`, `!=`)
2. `comparison` (`<`, `>`, `<=`, `>=`)
3. `term` (`+`, `-`)
4. `factor` (`*`, `/`)
5. `unary` (Prefix: `+`, `-`)
6. `primary` (Highest: `()`, `Integer`)

6.3 Resolving the "Dangling Else" Problem

A classic ambiguity in compiler design is the "Dangling Else".

- **Ambiguity:** `if (x) if (y) A; else B;`
- **Solution:** We relied on Bison's default `Shift` preference. When the parser encounters the `else`, it shifts it onto the stack, binding it to the nearest open `if`.

7 AST Structure & Construction

The Abstract Syntax Tree (AST) is the final output of the analysis phase.

7.1 Node Structure (`ast.h`)

We defined a generic `ASTNode` structure capable of representing any construct:

```
1 typedef struct ASTNode {  
2     NodeType type;           // NODE_BINOP, NODE_VAR_DECL, NODE_IF  
3     char* val_str;         // For Identifiers ("x") or Operators ("+")  
4     int val_int;            // For Integers (10)  
5     struct ASTNode *left;    // Left Child  
6     struct ASTNode *right;   // Right Child  
7 } ASTNode;
```

7.2 Bottom-Up Construction

The AST is built bottom-up using Bison actions. For example:

```
1 term: term PLUS factor {  
2     $$ = create_binop("+", $1, $3); // Connects Left ($1) and Right ($3)  
3 }
```

8 System Test Suite

To validate the compiler's robustness, we executed a comprehensive suite of **20 Test Cases** (10 Valid, 10 Invalid). The test script verified that valid programs return Exit Code 0 and invalid programs return Exit Code 1.

8.1 Valid Test Cases (Functional Verification)

#	Test Case	Description	Status
1	test01_basic	Simple variable declaration.	Pass
2	test02_math	Arithmetic precedence ($2 + 3 * 4$).	Pass
3	test03_ifelse	Basic branching logic.	Pass
4	test04_loop	while loop structure.	Pass
5	test05_nested	Stress Test (Loop inside Loop inside If).	Pass
6	test06_scope	Block scoping { var z = 5; }.	Pass
7	test07_compare	Relational operators (<, >).	Pass
8	test08_div_sub	Subtraction and Division logic.	Pass
9	test09_init_expr	Initialization with expressions.	Pass
10	test10_empty	Handling empty blocks and comments.	Pass

Table 1: Valid Test Cases

8.2 Invalid Test Cases (Error Handling)

#	Test Case	Description	Status
1	fail01_undeclared	Using variable without declaration.	Pass
2	fail02_double_decl	Declaring same variable twice.	Pass
3	fail03_no_semi	Missing semicolon (Syntax Error).	Pass
4	fail04_bad_assign	Assigning to a literal ($10 = x$).	Pass
5	fail05_paren	Mismatched parentheses.	Pass
6	fail06_brace	Unclosed block braces.	Pass
7	fail07_bad_char	Illegal input (var x = 10 \$).	Pass
8	fail08_keyword	Using keyword as identifier (var if).	Pass
9	fail09_missing_op	Incomplete expression (10 + ;).	Pass
10	fail10_bad_if	Malformed control flow syntax.	Pass

Table 2: Invalid Test Cases

9 Solutions & Design Decisions

This section summarizes critical technical solutions implemented during the lab.

1. Makefile Re-structuring:

- **Problem:** The `Makefile` was originally inside `src/`, complicating the build process from the root directory.
- **Solution:** We moved the `Makefile` to the root and used strict paths (e.g., `flex -o src/lex.yy.c src/lexer.l`), ensuring a clean separation of source and build artifacts.

2. Exit Code Management:

- **Problem:** The test script failed to detect errors because `main()` returned 0 even after `yyerror()` was called.
- **Solution:** We modified `yyerror()` to call `exit(1)`. This guarantees that the OS receives a failure signal, allowing the automated test suite to function correctly.

3. Symbol Table Safety:

- **Design:** Implemented a linear symbol table with a `MAX_VARS` limit of 1000.
- **Safety:** Added explicit bounds checking to prevent Segmentation Faults.

4. Runtime Safety Division by Zero :

A critical improvement was the addition of runtime error checking within the `eval()` function.

- **Problem:** Native C division (`x / 0`) causes a *Floating Point Exception* and crashes the compiler.
- **Solution:** We added an explicit check before any division operation in the AST evaluator:

```
if (strcmp(node->op, "/") == 0) {
    if (rhs == 0) {
        fprintf(stderr, "Runtime Error: Division by zero\n");
        exit(1);
    }
    return lhs / rhs;
}
```

This ensures the interpreter gracefully reports the error and exits, rather than crashing the system.

10 Limitations and Future Extensions

Although the current implementation fulfills all functional requirements for a basic compiler front-end, there are several areas where the design could be expanded.

10.1 Limitations

- **Symbol Table Efficiency:** The current symbol table is implemented as a linear array with a fixed limit (`MAX_VARS` 1000). For significantly larger programs, this would be inefficient ($O(n)$ lookup time). A production-grade compiler would use a Hash Table for $O(1)$ lookups.

- **Type System:** The language is strictly limited to the `integer` data type. It does not currently support floating-point numbers, strings, or boolean types.
- **Modular Programming:** The language does not support function definitions or function calls. The entire program is effectively a single `main` body.

10.2 Future Extensions

- **Type Checking:** The `check_symbol` function could be extended to store metadata about variable types, enabling the compiler to reject invalid assignments (e.g., assigning a string to an integer).
- **Code Generation:** The natural next step is to traverse the constructed AST and emit **LLVM IR** or **x86_64 Assembly**. This would transform the tool from a syntax checker into a fully functional compiler that produces executables.
- **Constant Folding:** We could implement an optimization pass during AST construction to pre-calculate constant expressions (e.g., turning `3 + 5` directly into `8`), reducing the runtime overhead of the final program.

11 Conclusion

This project successfully implements a front-end compiler satisfying all functional requirements. By integrating a recursive grammar with a robust AST construction strategy, the compiler handles complex control flows and mathematical expressions accurately. The decision to implement **Panic Mode** error handling ensures that invalid inputs are strictly rejected, providing a reliable foundation for future code generation stages.