

# Lab 6: Full-Stack Integration of a Managed Runtime Environment — Shell, Compiler, and Virtual Machine with Garbage Collection

Jithendra Gannavarapu (2025MCS2743)  
Manmadha Rao Kuttum (2025MCS2101)

February 3, 2026

## Abstract

This report details the successful integration of a multi-stage software toolchain. The system unifies a POSIX-compliant shell (Lab 1), a Flex/Bison front-end (Lab 2), a Bytecode IR Generator (Lab 3), and a managed Virtual Machine with an internal Mark-Sweep Garbage Collector (Labs 4 & 5). We demonstrate how metadata generated during parsing is utilized for instruction-level debugging and memory liveness analysis. The final system achieves full binary compatibility between independent modules through a strictly defined instruction set and value-tagging system.

## 1 Introduction

The objective of Lab 6 was to move beyond standalone implementations and create a unified "Integrated System." In this architecture, the shell is not just a command-line interface but a process coordinator. The compiler is not just a syntax checker but a bytecode emitter. The Virtual Machine (VM) is not just a mathematical evaluator but a state machine that supports introspection.

## 2 System Architecture and Codebase Organization

The system is organized into a modular pipeline where each directory represents a distinct phase of the software lifecycle.

- **01\_Shell:** The primary user interface. It manages program IDs (PIDs) and coordinates between the compiler and VM via the `ProgramManager`.
- **02\_Parser:** Contains the Lexer (`lexer.l`) and Parser (`parser.y`) developed in Lab 2 to transform source code into an Abstract Syntax Tree (AST).
- **03\_Compiler:** The IR Generator. It performs post-order traversal on the AST to emit a linear vector of bytecode opcodes.

- **04\_VM\_Execution:** The runtime engine (`VirtualMachine.cpp`) that executes bytecode using a Fetch-Decode-Execute cycle and supports software breakpoints.
- **05\_Memory\_GC:** The Mark-Sweep Garbage Collector developed in Lab 5 for dynamic memory reclamation and liveness tracking.

## 3 Formal Technical Specification

### 3.1 Instruction Set Architecture (ISA)

The system utilizes a custom 8-bit opcode format with 32-bit operands. Below are the primary instructions that unify the compiler and the VM:

Opcode	Mnemonic	Description
0x01	PUSH val	Pushes a 32-bit integer onto the operand stack.
0x10	ADD	Pops two values, adds them, and pushes the result.
0x21	JZ addr	Pops a value; if zero, sets PC to the target address.
0x30	STORE idx	Pops a value and stores it in the VM memory at index.
0x31	LOAD idx	Loads a value from memory index to the stack.
0x42	PRINT	Pops a value and displays it to the standard output.
0xFF	HALT	Terminates execution and returns control to the shell.

Table 1: Core Instruction Set for the Integrated System

### 3.2 Value Tagging and Object Representation

To support the Garbage Collector (Lab 5), every stack and memory element uses a `Value` tagged union.

$$Value = \{ \text{Type: } (INT|OBJ), \text{Data: } (int32\_t|Object*) \}$$

This allows the VM to distinguish between primitive data and heap-allocated objects during the **Mark Phase**.

## 4 Advanced Integration Features

### 4.1 Control Flow and Label Back-patching

One of the most complex integration points was the implementation of `WHILE` loops. Because jump targets (e.g., the end of a loop) are unknown when the jump instruction is first emitted, the `IRGenerator` uses a **Fixup List**.

1. The generator emits a jump instruction with a placeholder (0).
2. It records the bytecode offset of this placeholder in a fixup table.
3. Once the target label is "placed," it iterates through the fixup table and overwrites the placeholders with the actual memory address.

## 4.2 The Debugger State Machine

The Lab 6 debugger replaces the Lab 2 `ptrace` logic with internal state accessors. By modifying the VM to use an `executeNext()` method, the shell can execute precisely one instruction and then pause to display the "registers" (PC and Stack).

## 5 Garbage Collection & Memory Safety

The system utilizes the Mark-Sweep algorithm to handle transitive reachability. Unlike standalone versions, this GC is "Execution-Aware."

- **Mark Phase:** Recursively traces live objects starting from the current VM stack and global memory variables.
- **Sweep Phase:** Traverses the `objects` linked-list and unlinks nodes whose `marked` flag is false.

## 6 Integrated System Validation

### 6.1 Integration Case Study: Control Flow and Variable Bindings

We verified the system using a script containing an `if`-condition and a `while`-loop. The logs below demonstrate the data flow from AST to execution.

```
1 (debug) ast
2 WHILE
3   CONDITION: OPERATOR: < (ID: i, VALUE: 3)
4   BODY: ASSIGNMENT: i (OPERATOR: + (ID: i, VALUE: 1))
5
6 (debug) bytecode
7   25: 0x31 (LOAD ) 2
8   27: 0x01 (PUSH ) 3
9   29: 0x14 (CMP )
10  30: 0x21 (JZ    ) 41
11
12 (debug) regs
13 --- VM Register State ---
14 PC    : 2
15 Stack : [ 10 ]
```

Listing 1: Actual System Logs during Debugging

### 6.2 Analysis of Results

The trace above confirms that:

1. The **Parser** correctly interpreted the `WHILE` loop hierarchy.
2. The **Compiler** generated a `JZ` (Jump if Zero) at offset 30 pointing to the escape address 41.
3. The **VM** successfully stepped through instructions, showing the value 10 on the stack after the first assignment.

## 7 Conclusion

The Lab 6 Integrated System achieves the goal of a "Unified Execution Model." By centralizing control in a managed `ProgramManager` and implementing a step-enabled Virtual Machine, we have created a toolchain that is both executable and observable. This work demonstrates the critical necessity of metadata flow—where the front-end's syntax analysis directly informs the back-end's debugging and memory management capabilities.